

# LTAT.02.004 MACHINE LEARNING II

## **Performance evaluation**

Sven Laur  
University of Tartu

# Can we quantify model performance at all?



For some machine learning tasks there is no direct performance measure

- ▷ speech synthesis
- ▷ conversational chat bots
- ▷ image and music generation tasks

For all of these tasks we can define many surrogate measures

- ▷ understandability and expressiveness of speech
- ▷ grammatical correctness and adherence of other implicit patterns

# Why do we estimate performance?

- ▷ To estimate how does the algorithm perform in the future
  - ◇ This is the most important question in the practice
  - ◇ We are interested on performance of a particular predictor
- ▷ To find the best hyperparameter instance for our dataset
  - ◇ It is quite tricky task if we consider all subtleties
  - ◇ We are comparing different algorithm instances on our data
- ▷ To compare different algorithms and choose the best
  - ◇ This is needed to justify the development of a new algorithm
  - ◇ We are comparing average behaviour of algorithms
- ▷ To see if there is a dependence between input and the output
  - ◇ Studies in biology or sociology are all about causal dependencies
  - ◇ We are interested in statistically significant performance levels

## Short list of goodness measures

Some goodness measures for classification

- ▷ Accuracy – the percentage of correctly classified observations
- ▷ Precision – the percentage of correct labels among positive guesses
- ▷ Recall – the percentage of positive cases that are detected

Some goodness measures for regression

- ▷ Normalised mean square error
- ▷ Normalised mean absolute error
- ▷ Trimmed mean square and absolute error estimates

# Handwritten digit recognition task

MNIST database of handwritten digits

- ▷  $28 \times 28$  grayscale images of numbers
- ▷ Training set contains 60,000 images.
- ▷ Test set contains 10,000 images.
- ▷ We consider separation of two classes.



# Scaling laws

The performance of a classification algorithm depends on three factors:

- ▷ complexity of a model,
- ▷ the amount of training data,
- ▷ the amount of computational resources.

There are theoretical scaling laws.



- ◇ Statistical Learning Theory
- ◇ Gives conservative relation between dataset size and performance.

There are empirical scaling laws.

- ◇ Govern the architectural search of foundational models.
- ◇ Tells how much data and compute are needed to get target performance.

## Examples of scaling laws

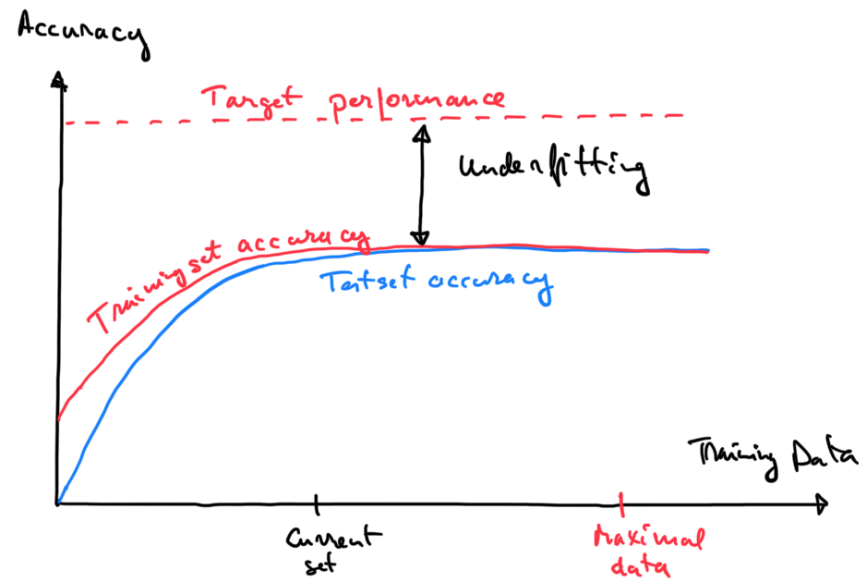
We can always fix the model and vary the compute or the dataset size.



There is always a discrepancy between test and training performance.

- ▷ This leaves some ambiguity about the limiting performance.
- ▷ Random choices are another source of ambiguity that limits utility.

# Data-bound learning tasks

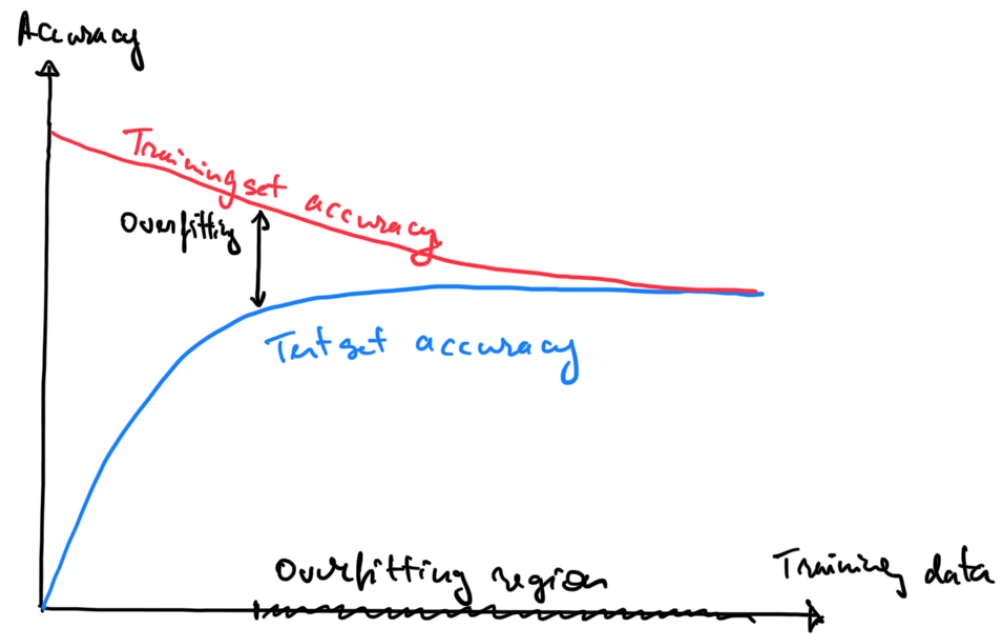


A learning task can be hard for following reasons:

- ▷ data bound – there is not enough (labelled) data
- ▷ algorithm bound – we do not have the right algorithm
- ▷ compute bound – we cannot finalise optimisation




# What is overfitting?




- ▷ Overfitting occurs when only the training performance decreases
- ▷ Overfitting = difference between training and limiting performance
- ▷ Overfitting is inevitable for stabilised models. The extent is important

# Models with distinct performance




- ▷ The MNIST test set consists of 10,000 samples
- ▷ Accuracy for reduced test sets consisting of 1,000 samples

# How much samples are needed?



- ▷ 9600 samples are needed to estimate accuracy with precision 1%
- ▷ If true accuracy  $\geq 90\%$  then the number of samples drops to 3500
- ▷ Performance increments of size 0.1 – 0.5% are relevant in practice
- ▷ This means test sets with sizes around 35,000 – 960,000

# Absolute vs relative performance




Relative difference in accuracy can be measured more precisely.

- ▷ Make predictions  $\hat{\mathbf{y}}_A$  and  $\hat{\mathbf{y}}_B$  on unlabelled data.
- ▷ Find data points on which predictions differ  $\mathcal{D} = \{i : \hat{\mathbf{y}}_A[i] \neq \hat{\mathbf{y}}_B[i]\}$ .
- ▷ Estimate relative difference in accuracies  $\Delta_{\mathcal{D}}$  on the set of differences  $\mathcal{D}$ .
- ▷ Estimate the relative size  $p_{\mathcal{D}}$  of  $\mathcal{D}$  and rescale the difference:

$$\text{accuracy}_A - \text{accuracy}_B = p_{\mathcal{D}} \cdot \Delta_{\mathcal{D}} .$$

## Shortcuts for relative performance



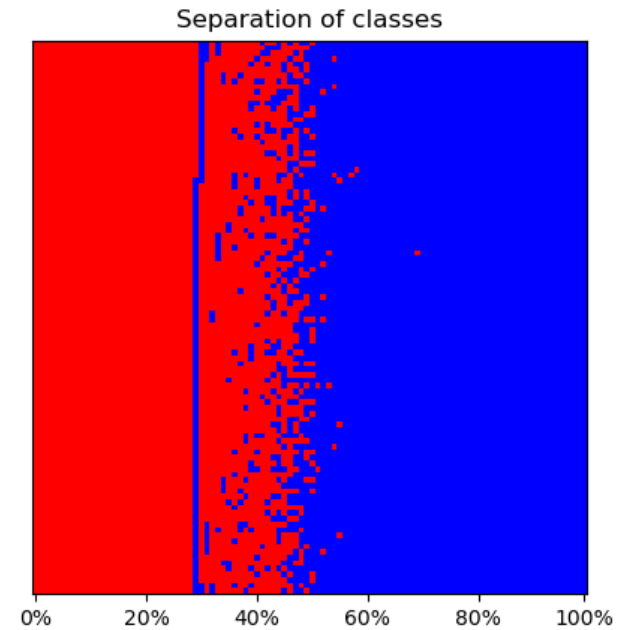

The alternative formula  $p_D \cdot \Delta_D$  is more approachable

- ▷ The proportion  $p_D$  can be computed without knowing true labels
- ▷  $\Delta_D$  can be estimated by labelling 1000 random samples from  $\mathcal{D}$

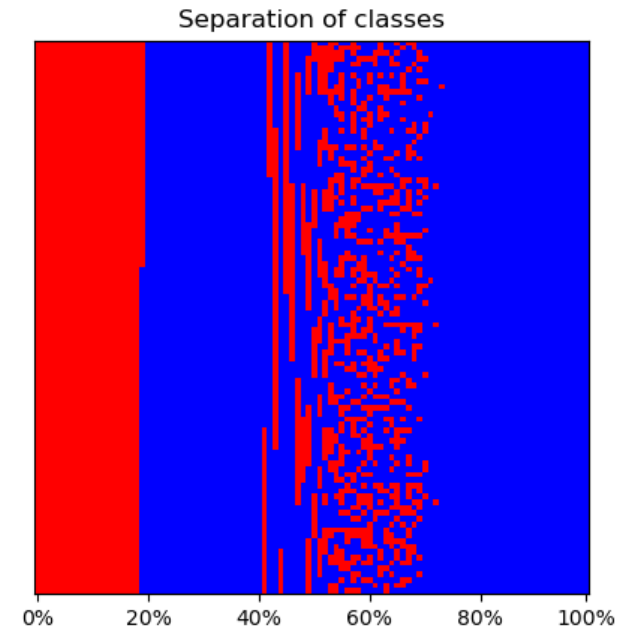

We can bound the relative difference without knowing true labels:

$$|\text{accuracy}_A - \text{accuracy}_B| \leq p$$

## Near ideal precision-recall graph

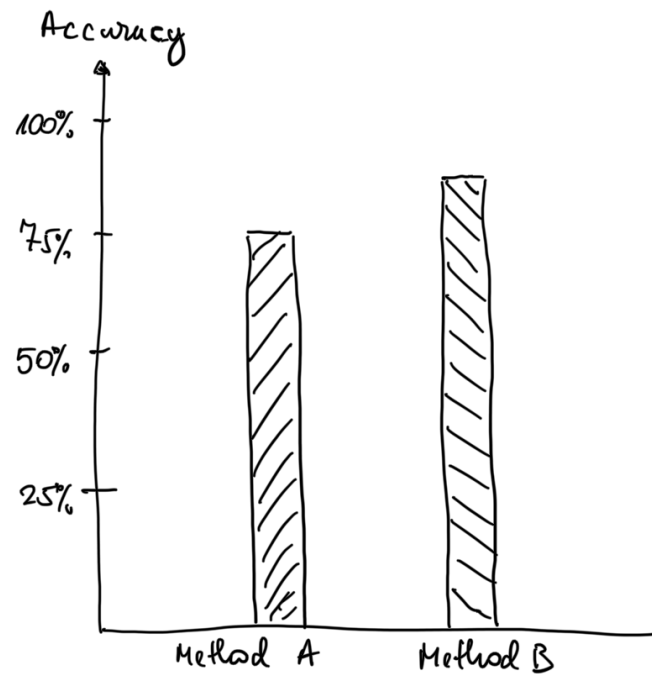


# Quick dropp-offs in precision-recall graph



- ▷ There must be a batch of negative data points with high scores.
- ▷ The score stabilises when the ratio of positives and negatives stabilises.

## Which model is better?



- ▷ Method A outputs the correct label in 75% cases on average
- ▷ Method B outputs the correct label in 85% cases on average
- ▷ Which method works significantly better on future data?



## How to model future?

```
accuracy = DataFrame(np.nan, index=range(1000), columns=['A', 'B'])
```

```
for i in range(1000):  
    data = generate_data(100)  
    accuracy.loc[i, 'A'] = evaluate_model_a(data)  
    accuracy.loc[i, 'B'] = evaluate_model_b(data)
```

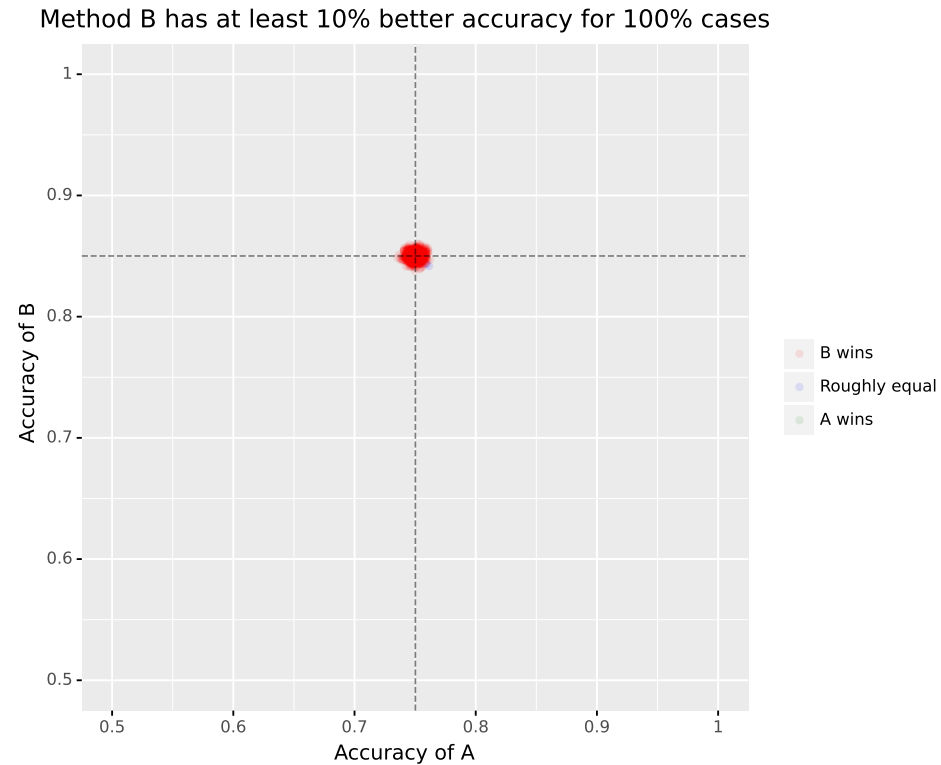
```
accuracy['status'] = 'Roughly equal'  
accuracy.loc[accuracy['A'] < 0.9 * accuracy['B'], 'status'] = 'B wins'  
accuracy.loc[accuracy['B'] < 0.9 * accuracy['A'], 'status'] = 'A wins'
```

# Future holds only 100 prediction tasks



Both methods are roughly equal although the accuracies are very different.

## Future holds 10000 prediction tasks



The method B is clearly superior to the method A as expected.

## Direct consequences

- ▷ The performance of a method fluctuates over all possible futures
- ▷ The magnitude of fluctuations decreases with the sample size
- ▷ True performance is defined as a limit over the infinite sample
- ▷ A finite estimate will always fluctuate around the true performance


Finite number of new observations in the future means that

- ▷ certain performance differences are irrelevant
- ▷ the pursuit of optimal performance is pointless

## Simplifying assumption

- ▷ We assume that the number of new observations is unbounded.

## How to estimate performance in the future



For any prediction algorithm we can find its expected goodness in the future.

**Practice.** Average goodness over a long enough series of future samples.

- ▷ Sampling should not change the data source in the future.
- ▷ All future samples should be independent from each other.

**Theory.** We should find expected goodness over the data distribution.

- ▷ The distribution always exists although we might not know it.
- ▷ Expected value exists even if the number of future samples is limited.

## Are these assumptions satisfied in practice?

### **Assumption I.** Data distribution does not change

- ▷ Some changes in data can be modelled
- ▷ If radical changes occur the model must be retrained
- ▷ Sometimes predictions must be valid regardless of inputs

### **Assumption II.** Future samples are independent from each other

- ▷ This assumption is always violated in text analysis
- ▷ This assumption is always violated in time-series analysis
- ▷ Correlation between future samples creates overconfidence
- ▷ This effect can be corrected with more careful sampling of a test set

# Notation and terminology

## Spaces

- ▷  $\mathcal{D}$  – data distribution
- ▷  $\mathcal{X}$  – input space, feature space
- ▷  $\mathcal{Y}$  – output space, target space
- ▷  $\mathcal{F} \subseteq \{f : \mathcal{X} \times \Omega \rightarrow \mathcal{Y}\}$  – model class

## Instances

- ▷  $\mathbf{x} \in \mathcal{X}$  – instance
- ▷  $y$  – true value of an instance, target value
- ▷  $\hat{y} = f(\mathbf{x})$  – predicted target value

## Loss:

- ▷  $L : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathcal{R}$  – the cost of using prediction  $\hat{y}$  instead of  $y$

## Theoretical formulation

Let  $\mathcal{D}$  be the distribution of  $(x, y)$  pairs where  $x$  is the input and  $y$  is the target of a prediction algorithm  $f$ .

Let  $L : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$  be the *loss function* which takes in the predicted value  $\hat{y}$  and the actual value  $y$  and outputs resulting loss.

Then the corresponding *risk*  $R(f)$  is computed as *mathematical expectation*

$$R(f) := \mathbf{E}_{\mathcal{D}}(L(f(x), y)) = \int_{(x,y) \in \mathcal{D}} L(f(x), y) dF(x, y)$$

where  $F$  is the corresponding probability measure.



## Practical example

- ▷ Let  $f(x_1, x_2) \equiv 0$  and let  $L(\hat{y}, y) = (y - \hat{y})^2$ . What is the risk  $R(f)$  if the next data sample is chosen uniformly from the following table.

$x_1$	$x_2$	$y$
0	0	0
0	0	1
0	1	1
1	0	0
1	1	1
0	0	0

- ▷ Propose a new prediction rule  $f_*$  that minimises the risk.
- ▷ Is there always a prediction rule that minimises the risk?

## Empirical risk estimation



When the sample  $D_N = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)\}$  is *representative* then we can approximate risk  $R(f)$  with *empirical risk*:

$$R_N(f) = \frac{1}{N} \cdot \sum_{i=1}^N L(f(\mathbf{x}_i), y_i) \ .$$

**IID sampling assumption.** The following conditions assure that the sample data  $D_N$  is representative (with high probability).

- ▷ All samples are independent from each other.
- ▷ All samples are drawn from the same distribution.
- ▷ Future samples come from the same distribution as the data  $D_N$ .

# Empirical risk



- ▷ depends on the dataset
- ▷ statistical fluctuations decrease with size

## Law of large numbers

**Central limit theorem.** Let  $z_1, \dots, z_N$  be independent and identically distributed samples from a *real-valued distribution* with a *finite standard deviation*  $\sigma$  and *mean*  $\mu$ . Then the random variable

$$S = \sqrt{N} \left( \frac{1}{N} \cdot \sum_{i=1}^N z_i - \mu \right)$$

converges *in distribution* to normal distribution  $\mathcal{N}(\text{mean} = 0, \text{sd} = \sigma)$ .

## Translation

Under mild assumptions the empirical risk  $R_N(f)$  converges to risk  $R(f)$  and we can actually use normal distribution to estimate probabilities:


$$\Pr [|R_N(f) - R(f)| \geq \varepsilon] \lesssim 2 \cdot \int_{-\infty}^{\varepsilon} \frac{\sqrt{N}}{\sqrt{2\pi}\sigma} \exp\left(-\frac{Nt^2}{2\sigma^2}\right) dt$$

for a finite value  $\sigma$  where  $\sigma^2$  is the variance of loss  $\mathbf{D}(R(f))$ .

## Reasoning

- ▷ If  $(\mathbf{x}_i, y_i)$  are IID samples then  $z_i = L(f(\mathbf{x}_i), y_i)$  are also IID samples.
- ▷ By definition  $\mu = \mathbf{E}(z) = \mathbf{E}(L(f(\mathbf{x}), y)) = R(f)$ .
- ▷ CLT assumes that risk  $\mu$  is finite and standard deviation  $\sigma$  is finite.

# Visual representation



Convergence implies that the centre area of is well approximated

- ▷ 90% confidence intervals are roughly the same for both distributions